

INTRO to DATA SCIENCE

LECTURE 11: DATABASES, STRUCTURED DATA, & INTRO TO SQL

Francesco Mosconi
DAT10 SF // November 8, 2014

RECAP

LAST TIME:

I. BIG DATA

II. PROGRAMMING MODEL

III. IMPLEMENTATION DETAILS

IV. WORD COUNT EXAMPLE

QUESTIONS?

AGENDA

LECTURE:

I. DATABASE EVOLUTION

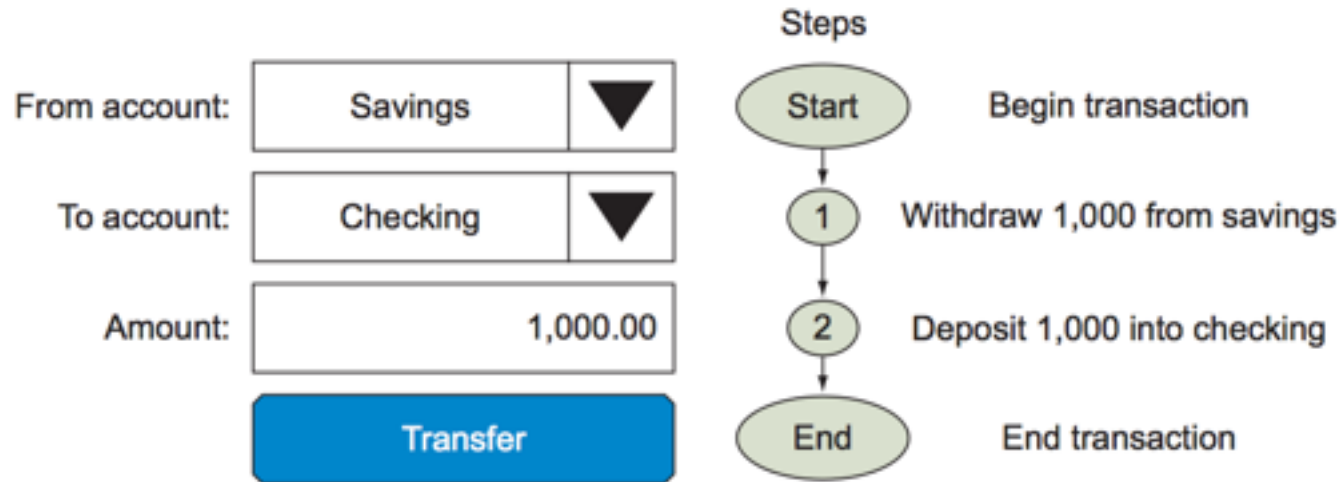
II. THE NOSQL MOVEMENT

III. WORKING WITH STRUCTURED DATA (MYSQL, SQLITE)

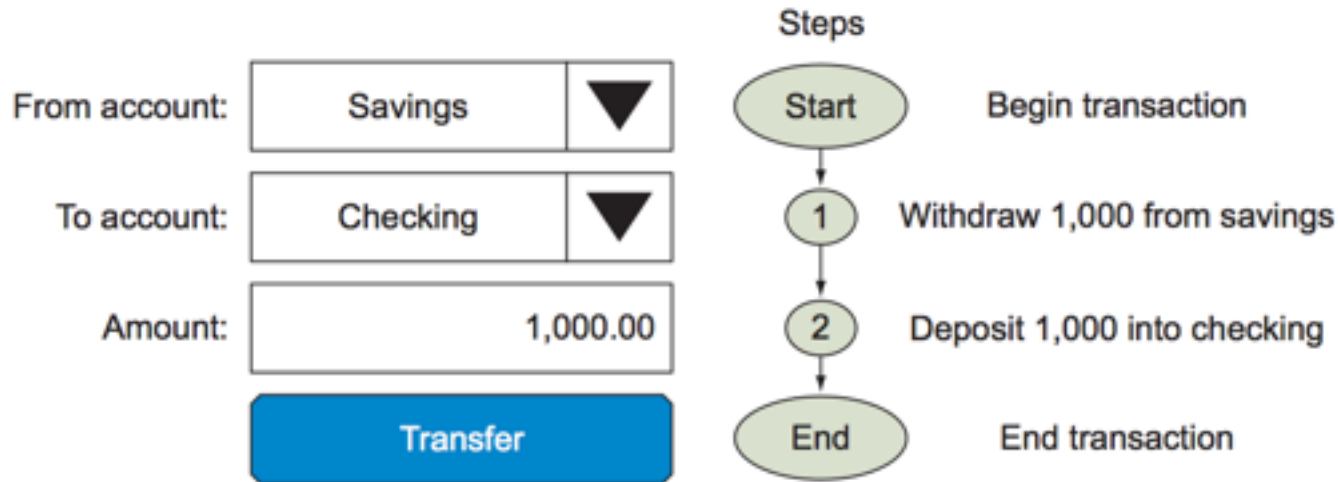
LAB: SQL (SQLITE)

I. THE EVOLUTION OF DATABASE TECHNOLOGY

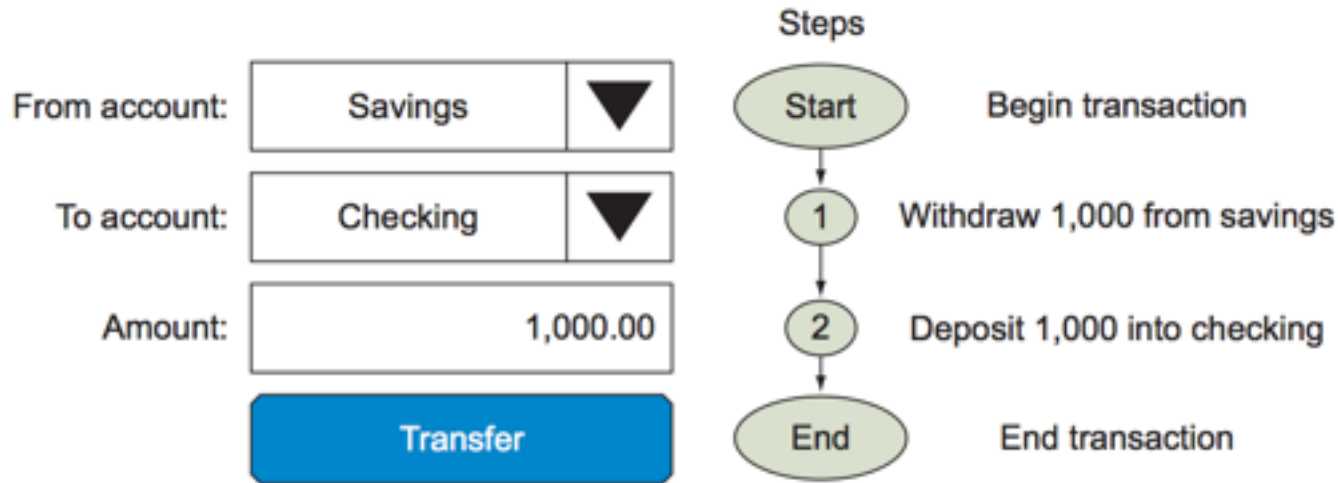
What is transactional integrity? A motivating example:



What happens if step 1 succeeds and step 2 fails?



What if you request your balance between step 1 and step2?



Transaction concepts:

- Transaction
- Begin / end transaction
- Rollback

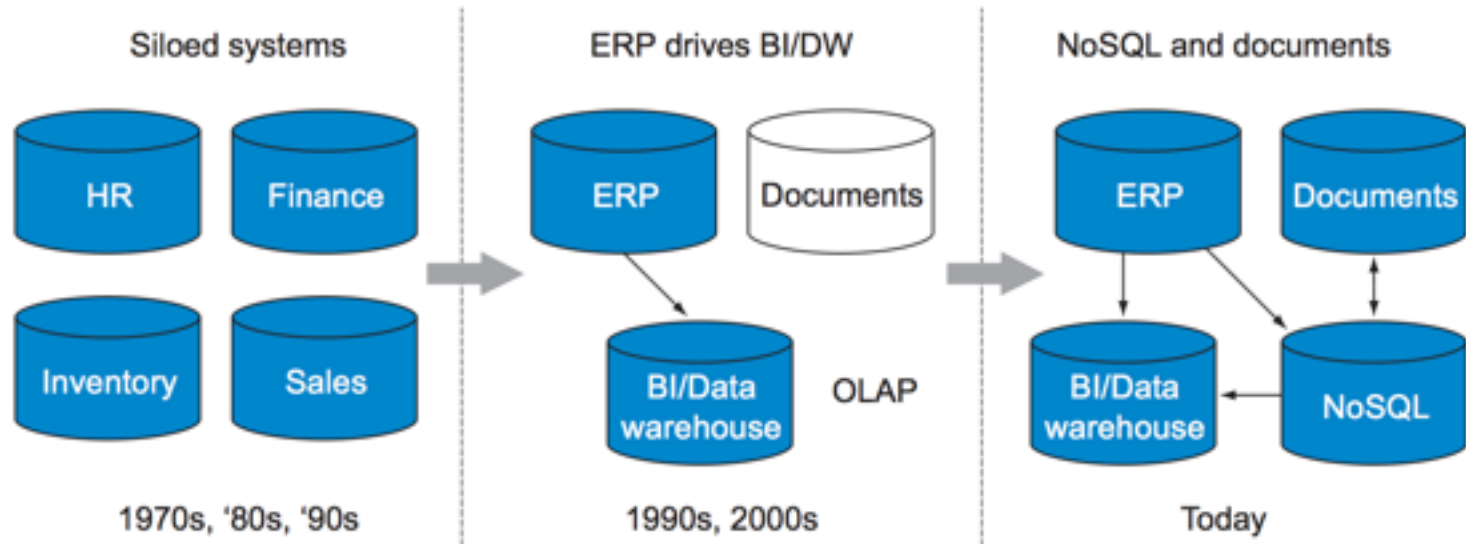
What other types of business activities can you think of that would be “transactions” as defined here...?

What other types of business activities can you think of that would be “transactions” as defined here...?

| | Debit | Credit |
|----------------|----------|----------|
| Asset | Increase | Decrease |
| Liability | Decrease | Increase |
| Income/Revenue | Decrease | Increase |
| Expense | Increase | Decrease |
| Equity/Capital | Decrease | Increase |

| | Account | Debit (Dr) | Credit (Cr) |
|----|------------|------------|-------------|
| 1. | Rent | 100 | |
| | Bank | | 100 |
| 2. | Bank | 50 | |
| | Sales | | 50 |
| 3. | Equipment | 5200 | |
| | Bank | | 5200 |
| 4. | Bank | 11000 | |
| | Loan | | 11000 |
| 5. | Salary | 5000 | |
| | Bank | | 5000 |
| 6. | Total (Dr) | 21350 | |
| | Total (Cr) | | 21350 |

That's why enterprise resource planning (ERP) systems and relational database management systems (RDBMS) grew up together.




- 1960s
 - Hierarchical data structure (IBM IMS)
 - Network data structure (CODASYL)
- 1970s
 - Relational data model
 - *A Relational Model of Data for Large Shared Data Banks* – E. F. Codd [1970]
 - System R (IBM), Ingres (Berkeley)

- 1980s
 - Commercialization of RDBMS
 - Oracle, Sybase, IBM DB2, Informix
 - SQL
 - ACID (Attomic, Consistent, Isolated, Durable)
- 1990s
 - PC RDBMS
 - Paradox, Microsoft SQL Server & Access
 - Larger DBs, driven by internet
 - Consolidation among commercial DB vendors

- 2000s
 - Commercialization of Open Source RDBMS
 - MySQL, Postgres
 - Evolving requirements expose RDBMS limitations
 - Storing complex and dynamic objects
 - Processing increasing data volumes
 - Analyzing massive amounts of data

Business drivers for NoSQL include:

- Volume
- Velocity
- Variability
- Agility



Business drivers for NoSQL include:

- Volume – the ability to query big data using clusters of commodity processors (horizontal scaling, parallel processing)
- Velocity – the ability to maintain performance in the face of traffic bursts from public-facing websites
- Variability – the ease of capturing & reporting on exception data
- Agility – object-relational mapping is complicated; even small changes can substantially slow development projects

II. THE NOSQL MOVEMENT

Eric Brewer's CAP (**C**onsistency, **A**vailability, **P**artition Tolerance) Theorem [2000]

For a distributed system -> Pick 2!

Research

MapReduce: Simplified Data Processing on Large Clusters – Google [2004]

Bigtable: A Distributed Storage System for Structured Data – Google [2006]

Dynamo: Amazon's Highly Available Key-value Store – Werner Vogels, et. al. [2007]

Pregel: A System for Large-Scale Graph Processing – Google [2010]

BASE (**B**asic **A**vailability, **S**oft-state, **E**ventually Consistent)

Vs.

Acid

- Get transaction details right
- Block any reports while you are working
- Be pessimistic: anything might go wrong!
- Detailed testing and failure mode analysis
- Lots of locks and unlocks



Base

- Never block a write
- Focus on throughput, not consistency
- Be optimistic: if one service fails it will eventually get caught up
- Some reports may be inconsistent for a while, but don't worry
- Keep things simple and avoid locks



| Type | Typical usage | Examples |
|---|--|--|
| <i>Key-value store</i> —A simple data storage system that uses a key to access a value | <ul style="list-style-type: none">• Image stores• Key-based filesystems• Object cache• Systems designed to scale | <ul style="list-style-type: none">• Berkeley DB• Memcache• Redis• Riak• DynamoDB |
| <i>Column family store</i> —A sparse matrix system that uses a row and a column as keys | <ul style="list-style-type: none">• Web crawler results• Big data problems that can relax consistency rules | <ul style="list-style-type: none">• Apache HBase• Apache Cassandra• Hypertable• Apache Accumulo |
| <i>Graph store</i> —For relationship-intensive problems | <ul style="list-style-type: none">• Social networks• Fraud detection• Relationship-heavy data | <ul style="list-style-type: none">• Neo4j• AllegroGraph• Bigdata (RDF data store)• InfiniteGraph (Objectivity) |
| <i>Document store</i> —Storing hierarchical data structures directly in the database | <ul style="list-style-type: none">• High-variability data• Document search• Integration hubs• Web content management• Publishing | <ul style="list-style-type: none">• MongoDB (10Gen)• CouchDB• Couchbase• MarkLogic• eXist-db• Berkeley DB XML |

Key-value

memcached, Redis, Riak, Tokyo Cabinet, Voldemort, Amazon SimpleDB

Column-oriented (Bigtable clones)

Cassandra, HBase

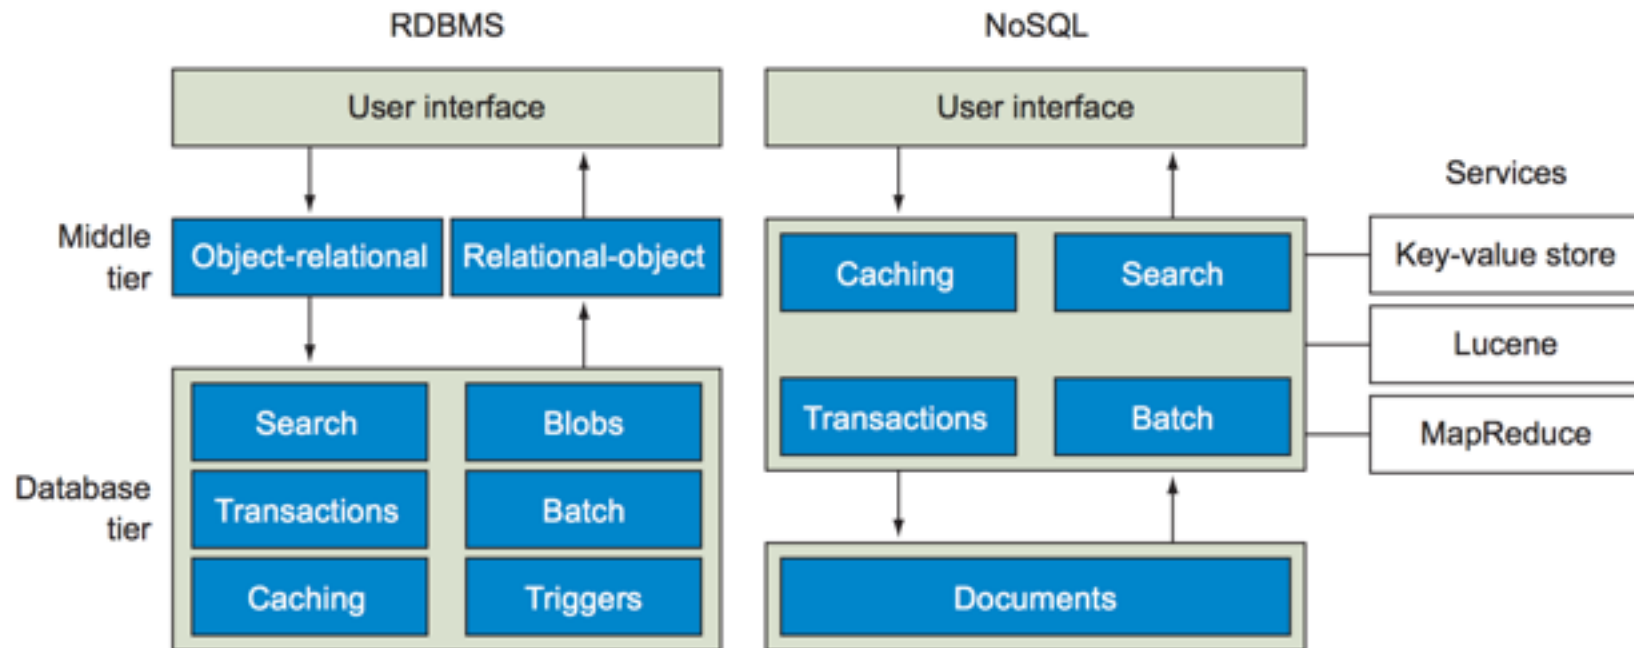
Document-oriented


MongoDB, CouchDB

Graph

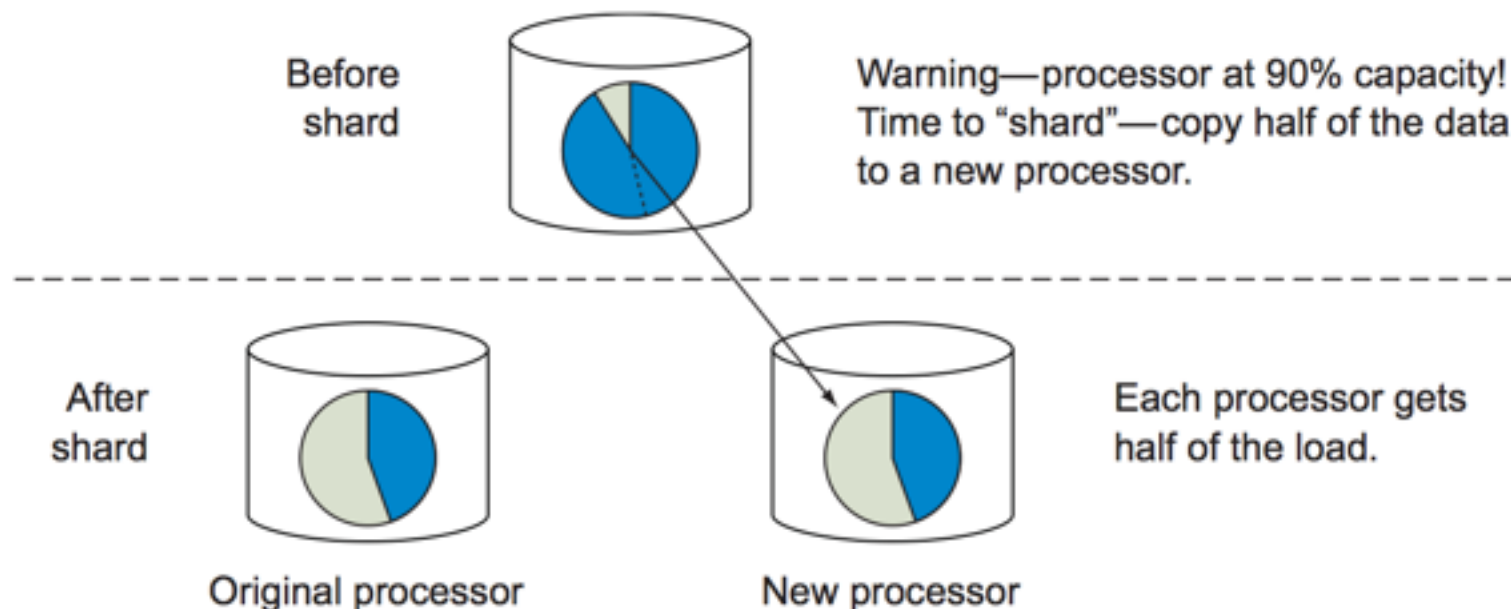
Neo4J, FlockDB, OrientDB, Pregel (Google)

| Case study/standard | Driver | Finding |
|------------------------|---|---|
| LiveJournal's Memcache | Need to increase performance of database queries. | By using hashing and caching, data in RAM can be shared. This cuts down the number of read requests sent to the database, increasing performance. |
| Google's MapReduce | Need to index billions of web pages for search using low-cost hardware. | By using parallel processing, indexing billions of web pages can be done quickly with a large number of commodity processors. |
| Google's Bigtable | Need to flexibly store tabular data in a distributed system. | By using a sparse matrix approach, users can think of all data as being stored in a single table with billions of rows and millions of columns without the need for up-front data modeling. |
| Amazon's Dynamo | Need to accept a web order 24 hours a day, 7 days a week. | A key-value store with a simple interface can be replicated even when there are large volumes of data to be processed. |
| MarkLogic | Need to query large collections of XML documents stored on commodity hardware using standard query languages. | By distributing queries to commodity servers that contain indexes of XML documents, each server can be responsible for processing data in its own local disk and returning the results to a query server. |

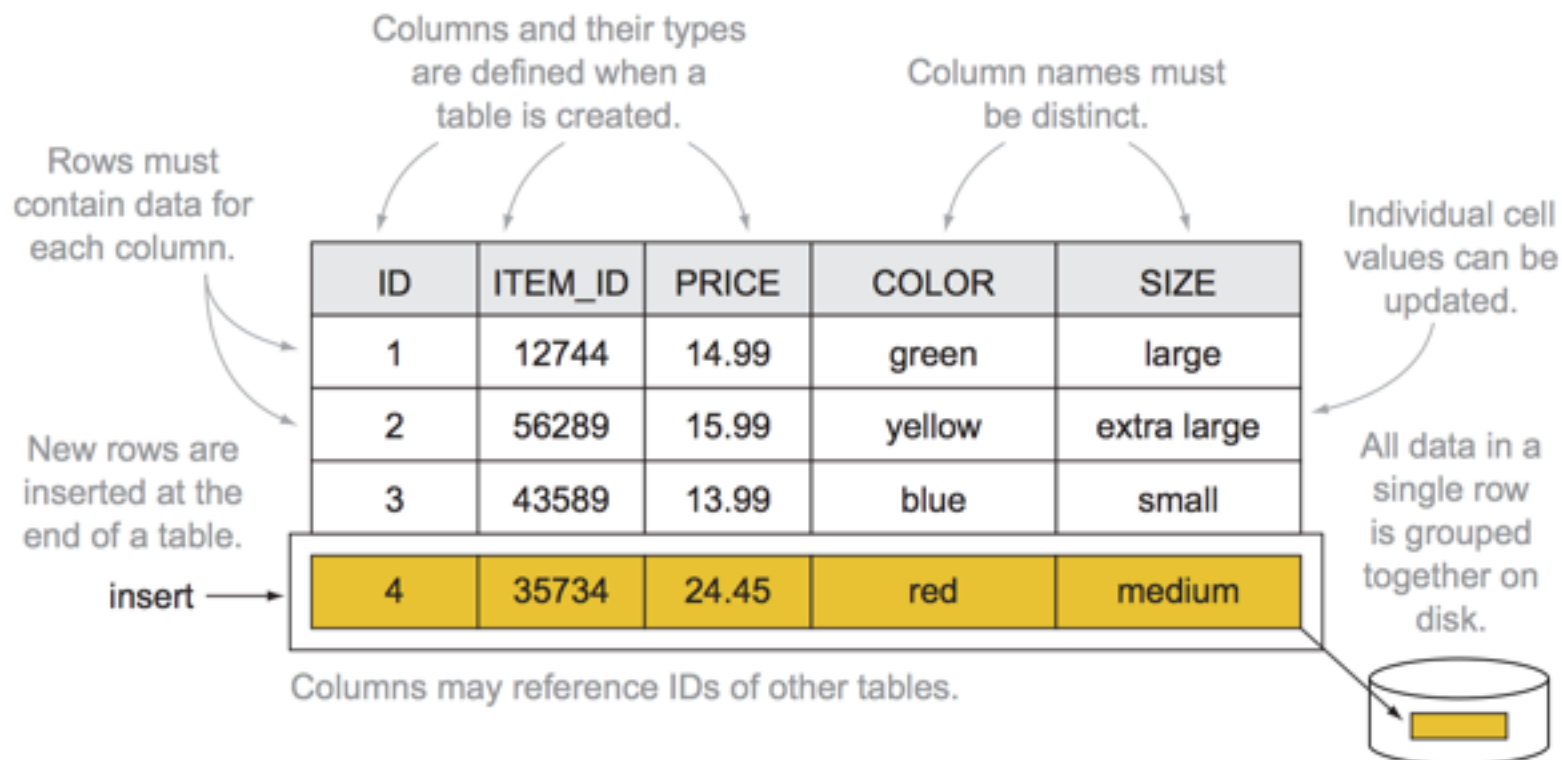


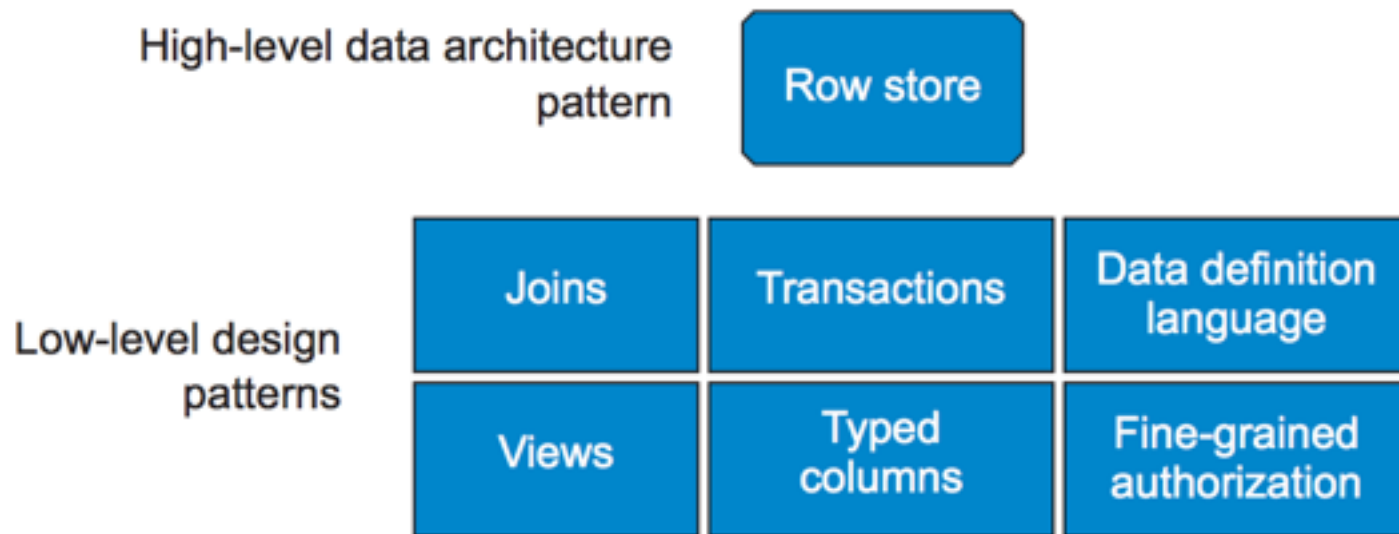


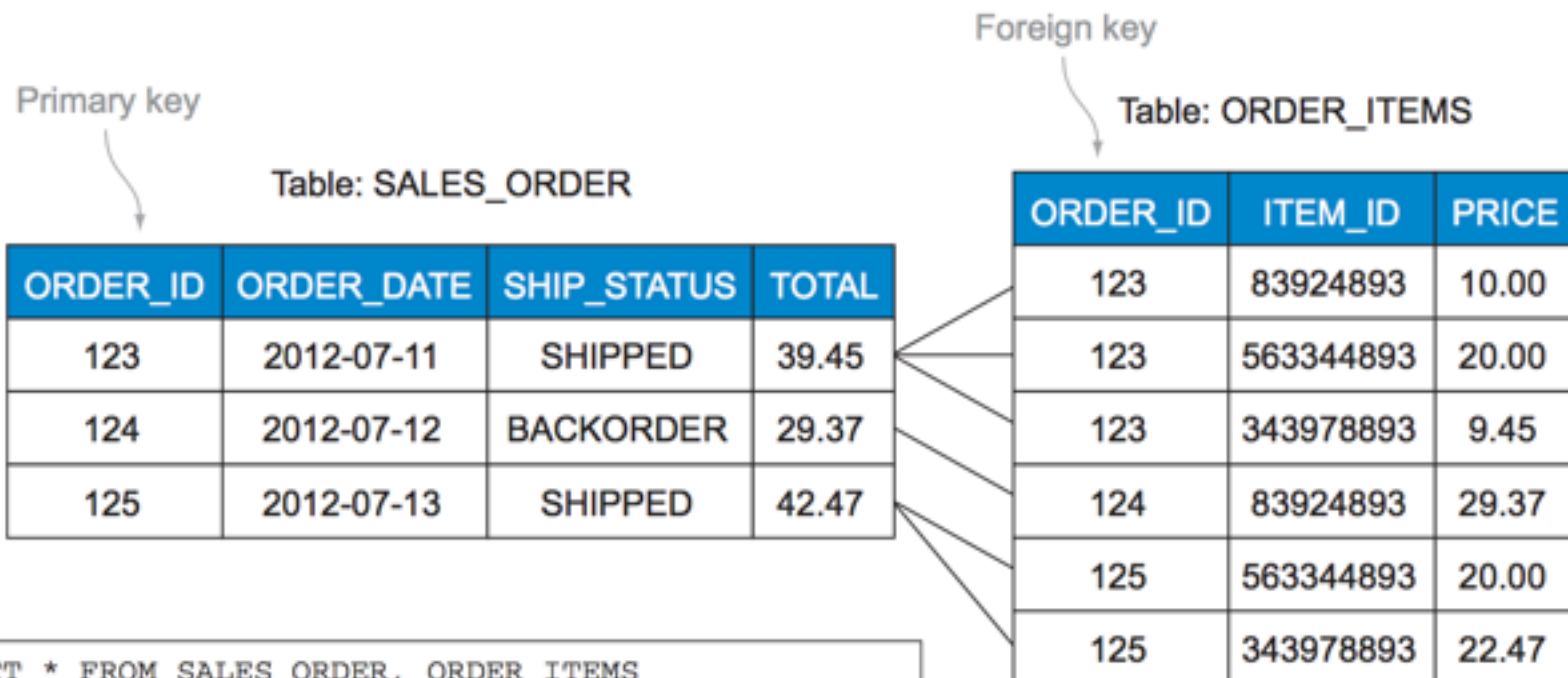
```
let $hash := hash($invoice, 'md5')
```

III. WORKING WITH STRUCTURED DATA (MYSQL, SQLITE)







```
SELECT * FROM SALES_ORDER, ORDER_ITEMS
WHERE SALES_ORDER.ORDER_ID = ORDER_ITEMS.ORDER_ID
```

We want to restrict
general access to this
column.

Physical table

| ORDER_ID | ORDER_DATE | SHIP_STATUS | CARD_INFO | TOTAL |
|----------|------------|-------------|--------------|-------|
| 123 | 2012-07-11 | SHIPPED | VISA-1234... | 39.45 |
| 124 | 2012-07-12 | BACKORDER | MC-5678... | 29.37 |
| 125 | 2012-07-13 | SHIPPED | AMEX-9012... | 42.47 |

The physical table includes all the column, including credit card info. Only select users ever see the physical table.

View of table

| ORDER_ID | ORDER_DATE | SHIP_STATUS | TOTAL |
|----------|------------|-------------|-------|
| 123 | 2012-07-11 | SHIPPED | 39.45 |
| 124 | 2012-07-12 | BACKORDER | 29.37 |
| 125 | 2012-07-13 | SHIPPED | 42.47 |

The view excludes some fields like credit card information. All sales analysts have access to the views.

| Feature | Strength | Weakness |
|--|--|---|
| Joins between tables | New views of data from different tables can easily be created. | All tables must be on the same server to make joins run efficiently. This makes it difficult to scale to more than one processor. |
| Transactions | Defining begin point, end point, and completion of critical transactions in an application is simple. | Read and write transactions may be slowed during critical times in a transaction unless the transaction isolation level is changed. |
| Fixed data definitions and typed columns | Easy way to define structure and enforce business rules when tables are created. You can verify on insert that all data conforms to specific rules. Allows range indexes over columns. | Difficult to work with highly variable and exception data when adding to a column. |
| Fine-grained security | Data access control by row and column can be done with a series of view and grant statements. | Setup and testing security access for many roles can be a complex process. |
| Document integration | None. Few RDBMSs are designed to easily query document structures. | Difficult to create reports using both structured and unstructured data. |

INTRO TO DATA SCIENCE

LAB: SQL